

Loan Eligibility Prediction dataset

January 28, 2025

1 Loan Eligibility Prediction

```
[2]: ## Importing Necessary Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
[3]: ## Read the dataset
df = pd.read_csv('loan_prediction_dataset.csv')
df.head()
```

```
[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

2 EDA

```
[4]: df.shape
```

```
[4]: (614, 13)
```

```
[5]: df.columns
```

```
[5]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',  
        'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
        'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],  
        dtype='object')
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 614 entries, 0 to 613  
Data columns (total 13 columns):  
#   Column                Non-Null Count  Dtype    
---  ---                     -  
0   Loan_ID               614 non-null   object   
1   Gender                601 non-null   object   
2   Married               611 non-null   object   
3   Dependents            599 non-null   object   
4   Education             614 non-null   object   
5   Self_Employed         582 non-null   object   
6   ApplicantIncome       614 non-null   int64    
7   CoapplicantIncome     614 non-null   float64  
8   LoanAmount            592 non-null   float64  
9   Loan_Amount_Term      600 non-null   float64  
10  Credit_History        564 non-null   float64  
11  Property_Area         614 non-null   object   
12  Loan_Status           614 non-null   object   
dtypes: float64(4), int64(1), object(8)  
memory usage: 62.5+ KB
```


```
[7]: df.isnull().sum()
```

```
[7]: Loan_ID           0  
    Gender          13  
    Married         3  
    Dependents      15  
    Education        0  
    Self_Employed   32  
    ApplicantIncome  0  
    CoapplicantIncome 0  
    LoanAmount      22  
    Loan_Amount_Term 14
```

```
Credit_History      50
Property_Area       0
Loan_Status         0
dtype: int64
```

```
[8]: ## Check for any outliers
plt.figure(figsize =(12,8))
sns.boxplot(data = df)
```

[8]: <Axes: >



```
[9]: ## Fill the null values of numerical datatype
df['LoanAmount'] = df['LoanAmount'].fillna(df['LoanAmount'].median())
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].
↳mean())
df['Credit_History'] = df['Credit_History'].fillna(df['Credit_History'].mean())
```

```
[10]: ## Fill the null values of Object datatype
df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])
df['Married'] = df['Married'].fillna(df['Married'].mode()[0])
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])
df['Self_Employed'] = df['Self_Employed'].fillna(df['Self_Employed'].mode()[0])
```


```
[11]: df.isnull().sum()
```

```
[11]: Loan_ID      0
      Gender      0
      Married     0
      Dependents  0
      Education   0
      Self_Employed 0
      ApplicantIncome 0
      CoapplicantIncome 0
      LoanAmount   0
      Loan_Amount_Term 0
      Credit_History 0
      Property_Area 0
      Loan_Status  0
      dtype: int64
```

```
[12]: ## EDA
      ## Number of people who took loan by gender
      print('Number of people who took loan by Gender')
      print(df['Gender'].value_counts())
      sns.countplot(x='Gender',data = df,palette='Set1')
```

```
Number of people who took loan by Gender
Gender
Male      502
Female    112
Name: count, dtype: int64
```

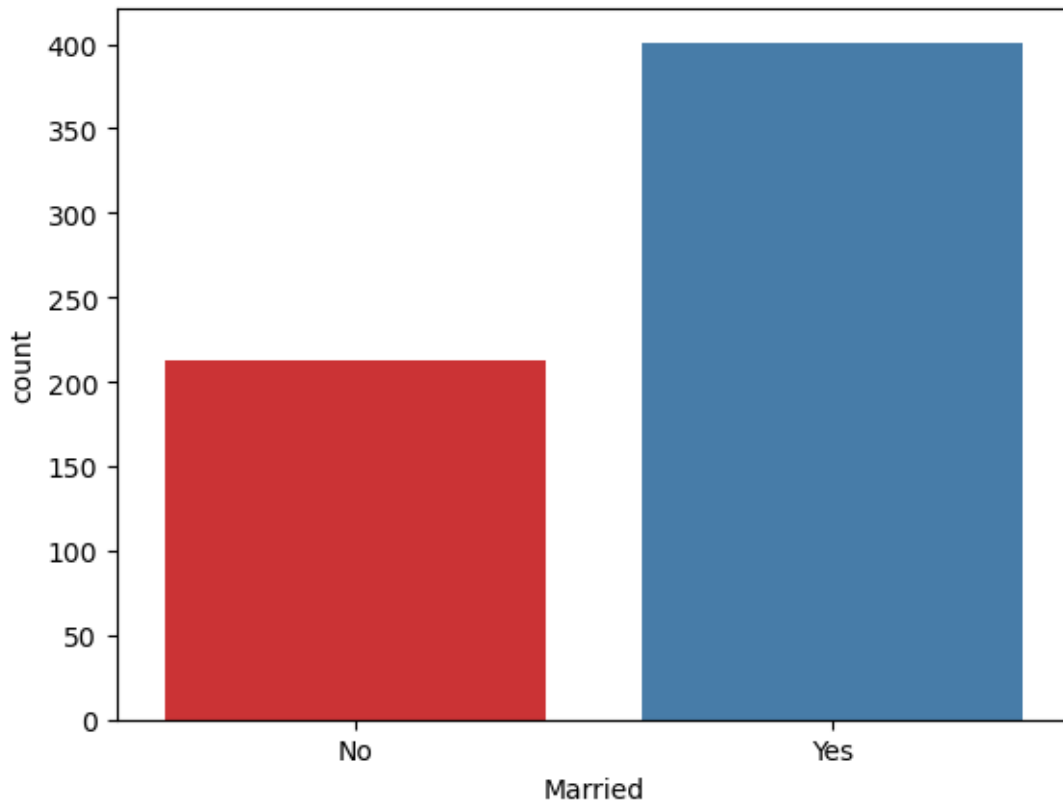
```
[12]: <Axes: xlabel='Gender', ylabel='count'>
```



```
[13]: ## Number of people who took loan by Martial status
print('Number of people who took loan by Martial status')
print(df['Married'].value_counts())
sns.countplot(x='Married',data = df,palette='Set1')
```

```
Number of people who took loan by Martial status
Married
Yes    401
No     213
Name: count, dtype: int64
```


```
[13]: <Axes: xlabel='Married', ylabel='count'>
```



```
[14]: ## Number of people who took loan by Education
print('Number of people who took loan by Education')
print(df['Education'].value_counts())
sns.countplot(x='Education',data = df,palette='Set1')
```

```
Number of people who took loan by Education
Education
Graduate      480
Not Graduate   134
Name: count, dtype: int64
```

```
[14]: <Axes: xlabel='Education', ylabel='count'>
```



```
[15]: ## Correlation check
numeric_df = df.select_dtypes(include=[np.number])
corr = numeric_df.corr()
print(corr)
plt.figure(figsize=(10,8))
sns.heatmap(corr,annot= True,cmap='BuPu')
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	\
ApplicantIncome	1.000000	-0.116605	0.565181	
CoapplicantIncome	-0.116605	1.000000	0.189218	
LoanAmount	0.565181	0.189218	1.000000	
Loan_Amount_Term	-0.045242	-0.059675	0.039235	
Credit_History	-0.014477	-0.001665	-0.007031	

	Loan_Amount_Term	Credit_History
ApplicantIncome	-0.045242	-0.014477
CoapplicantIncome	-0.059675	-0.001665
LoanAmount	0.039235	-0.007031
Loan_Amount_Term	1.000000	0.001395
Credit_History	0.001395	1.000000

[15]: <Axes: >



3 Feature Enginerring

```
[16]: ## Feature Enginerring
      ### Total Applicant Income
      df['Total_Applicant_Income'] = df['ApplicantIncome'] + df['CoapplicantIncome']
      df.head()
```

```
[16]:   Loan_ID Gender Married Dependents Education Self_Employed \
0  LP001002  Male    No         0    Graduate           No
1  LP001003  Male   Yes         1    Graduate           No
2  LP001005  Male   Yes         0    Graduate           Yes
3  LP001006  Male   Yes         0  Not Graduate           No
4  LP001008  Male   No         0    Graduate           No
```



	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
0	5849	0.0	128.0	360.0
1	4583	1508.0	128.0	360.0
2	3000	0.0	66.0	360.0
3	2583	2358.0	120.0	360.0
4	6000	0.0	141.0	360.0

	Credit_History	Property_Area	Loan_Status	Total_Applicant_Income
0	1.0	Urban	Y	5849.0
1	1.0	Rural	N	6091.0
2	1.0	Urban	Y	3000.0
3	1.0	Urban	Y	4941.0
4	1.0	Urban	Y	6000.0

```
[17]: ## Apply Log Trnasformation for ApplicantIncome


df['ApplicantIncomelog'] = np.log(df['ApplicantIncome'] + 1)
sns.distplot(df['ApplicantIncomelog'])
```

```
[17]: <Axes: xlabel='ApplicantIncomelog', ylabel='Density'>
```




```
[18]: ## Apply Log Trnasformation for LoanAmount  
  
df['LoanAmountlog'] = np.log(df['LoanAmount'] + 1)  
sns.distplot(df['LoanAmountlog'])
```

```
[18]: <Axes: xlabel='LoanAmountlog', ylabel='Density'>
```




```
[19]: ## Apply Log Trnasformation for Loan_Amount_Term  
  
df['LoanAmount_Term_log'] = np.log(df['Loan_Amount_Term'] + 1)  
sns.distplot(df['LoanAmount_Term_log'])
```

```
[19]: <Axes: xlabel='LoanAmount_Term_log', ylabel='Density'>
```



```
[20]: ## Apply Log Trnasformation for Total_Applicant_Income  
  
df['Total_Applicant_Income_Log'] = np.log(df['Total_Applicant_Income'] + 1)  
sns.distplot(df['Total_Applicant_Income_Log'])
```

```
[20]: <Axes: xlabel='Total_Applicant_Income_Log', ylabel='Density'>
```



```
[21]: ## DataFrame after applying log function
df.head()
```

```
[21]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	128.0	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status	Total_Applicant_Income	\
0	1.0	Urban	Y	5849.0	
1	1.0	Rural	N	6091.0	
2	1.0	Urban	Y	3000.0	

3	1.0	Urban	Y	4941.0
4	1.0	Urban	Y	6000.0

	ApplicantIncomelog	LoanAmountlog	LoanAmount_Term_log \
0	8.674197	4.859812	5.888878
1	8.430327	4.859812	5.888878
2	8.006701	4.204693	5.888878
3	7.857094	4.795791	5.888878
4	8.699681	4.955827	5.888878

	Total_Applicant_Income_Log
0	8.674197
1	8.714732
2	8.006701
3	8.505525
4	8.699681

```
[22]: ## Drop Unnecessary Features
cols = [
    'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Total_Applicant_Inc
df = df.drop(columns = cols, axis = 1)
df.head()
```

	Gender	Married	Dependents	Education	Self_Employed	Credit_History \
0	Male	No	0	Graduate	No	1.0
1	Male	Yes	1	Graduate	No	1.0
2	Male	Yes	0	Graduate	Yes	1.0
3	Male	Yes	0	Not Graduate	No	1.0
4	Male	No	0	Graduate	No	1.0

	Property_Area	Loan_Status	ApplicantIncomelog	LoanAmountlog \
0	Urban	Y	8.674197	4.859812
1	Rural	N	8.430327	4.859812
2	Urban	Y	8.006701	4.204693
3	Urban	Y	7.857094	4.795791
4	Urban	Y	8.699681	4.955827

	LoanAmount_Term_log	Total_Applicant_Income_Log
0	5.888878	8.674197
1	5.888878	8.714732
2	5.888878	8.006701
3	5.888878	8.505525
4	5.888878	8.699681

```
[23]: ## Converting all categorical to numerical do encoding technique
from sklearn.preprocessing import LabelEncoder
```

```
cols = [
    'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']
le = LabelEncoder()
for col in cols:
    df[col] = le.fit_transform(df[col])
```

```
[24]: df.head()
```

```
[24]:
```

	Gender	Married	Dependents	Education	Self_Employed	Credit_History	\
0	1	0	0	0	0	1.0	
1	1	1	1	0	0	1.0	
2	1	1	0	0	1	1.0	
3	1	1	0	1	0	1.0	
4	1	0	0	0	0	1.0	

	Property_Area	Loan_Status	ApplicantIncomeLog	LoanAmountLog	\
0	2	1	8.674197	4.859812	
1	0	0	8.430327	4.859812	
2	2	1	8.006701	4.204693	
3	2	1	7.857094	4.795791	
4	2	1	8.699681	4.955827	

	LoanAmount_Term_log	Total_Applicant_Income_Log
0	5.888878	8.674197
1	5.888878	8.714732
2	5.888878	8.006701
3	5.888878	8.505525
4	5.888878	8.699681

```
[25]: df.dtypes
```

```
[25]: Gender                int32
Married                    int32
Dependents                 int32
Education                  int32
Self_Employed              int32
Credit_History             float64
Property_Area              int32
Loan_Status                int32
ApplicantIncomeLog         float64
LoanAmountLog              float64
LoanAmount_Term_log        float64
Total_Applicant_Income_Log float64
dtype: object
```

```
[26]: ## Splitting Independent and Dependent features
```

```
X = df.drop(columns = ['Loan_Status'], axis = 1)
Y = df['Loan_Status']
```

4 Implementing Models

```
[27]: ## Modelling Part
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score , confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
[28]: ## Splitting Train and Test Data
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.
↳25,random_state= 42)
```

```
[30]: ## First Model 1 : Logistic Regression

model1 = LogisticRegression()
model1.fit(X_train,Y_train)
y_pred_model1 = model1.predict(X_test)
accuracy = accuracy_score(Y_test,y_pred_model1)
print("Accuracy Score of Logistic Regression: " ,accuracy * 100)
```

Accuracy Score of Logistic Regression: 77.27272727272727

```
[31]: score = cross_val_score(model1,X,Y,cv = 5)
score
print ("Cross Validation Score for Logistic Regression: " , np.mean(score)*100)
```

Cross Validation Score for Logistic Regression: 80.9462881514061

```
[ ]:
```

```
[32]: ## Second Model 2 : Decision Tree

model2 = DecisionTreeClassifier()
model2.fit(X_train,Y_train)
y_pred_model2 = model2.predict(X_test)
accuracy = accuracy_score(Y_test,y_pred_model2)
print("Accuracy Score of Decison Tree : " ,accuracy * 100)
```

Accuracy Score of Decison Tree : 70.12987012987013

```
[33]: score = cross_val_score(model2,X,Y,cv = 5)
score
print ("Cross Validation Score for Decision Tree " , np.mean(score)*100)
```

Cross Validation Score for Decision Tree 70.52245768359323

[]:

[34]: *## Third Model 3 : Random Classifier*

```
model3 = RandomForestClassifier()
model3.fit(X_train,Y_train)
y_pred_model3 = model3.predict(X_test)
accuracy = accuracy_score(y_pred_model3,Y_test)
print("Accuracy Score of Random Forest: " ,accuracy * 100)
```

Accuracy Score of Random Forest: 76.62337662337663

[]:

[35]: *## Fourth Model 4 : KNeighbor Model*

```
model4 = KNeighborsClassifier(n_neighbors = 3)
model4.fit(X_train,Y_train)
y_pred_model4 = model4.predict(X_test)
accuracy = accuracy_score(y_pred_model4,Y_test)
print("Accuracy Score of KNeighbor Model: " ,accuracy * 100)
```

Accuracy Score of KNeighbor Model: 71.42857142857143

[36]: score = cross_val_score(model4,X,Y,cv = 5)

```
score
print ("Cross Validation Score for KNeighbor " , np.mean(score)*100)
```

Cross Validation Score for KNeighbor 73.61721977875517

[37]: *## Checking the reports of all models*

```
from sklearn.metrics import classification_report

def generate_classification_report(model_name, Y_test, y_pred):
    report = classification_report(Y_test, y_pred)
    print(f"Classification Report for {model_name}:\n{report}\n")

generate_classification_report(model1, Y_test, y_pred_model1)
generate_classification_report(model2, Y_test, y_pred_model2)
generate_classification_report(model3, Y_test, y_pred_model3)
generate_classification_report(model4, Y_test, y_pred_model4)
```

Classification Report for LogisticRegression():

	precision	recall	f1-score	support
0	0.91	0.39	0.55	54
1	0.75	0.98	0.85	100

accuracy			0.77	154
macro avg	0.83	0.68	0.70	154
weighted avg	0.81	0.77	0.74	154

Classification Report for DecisionTreeClassifier():

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.59	0.48	0.53	54
1	0.75	0.82	0.78	100

accuracy			0.70	154
macro avg	0.67	0.65	0.66	154
weighted avg	0.69	0.70	0.69	154

Classification Report for RandomForestClassifier():

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.82	0.43	0.56	54
1	0.75	0.95	0.84	100

accuracy			0.77	154
macro avg	0.79	0.69	0.70	154
weighted avg	0.78	0.77	0.74	154

Classification Report for KNeighborsClassifier(n_neighbors=3):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.63	0.44	0.52	54
1	0.74	0.86	0.80	100

accuracy			0.71	154
macro avg	0.69	0.65	0.66	154
weighted avg	0.70	0.71	0.70	154

```
[38]: df['Loan_Status'].value_counts()
```

```
[38]: Loan_Status
1      422
0      192
Name: count, dtype: int64
```

```
[39]: ## Doing Oversampling for Imbalanced Data
```

```
from imblearn.over_sampling import RandomOverSampler
```

```
[40]: oversample = RandomOverSampler(random_state = 42)
X_resampled,Y_resampled = oversample.fit_resample(X,Y)

df_resampled = pd.concat([pd.DataFrame(X_resampled , columns = X.columns),pd.
↳Series(Y_resampled,name= "Loan_status")],axis=1)
```

```
[41]: X_resampled
```

```
[41]:      Gender  Married  Dependents  Education  Self_Employed  Credit_History  \
0         1         0           0           0           0           1.000000
1         1         1           1           0           0           1.000000
2         1         1           0           0           1           1.000000
3         1         1           0           1           0           1.000000
4         1         0           0           0           0           1.000000
..      ...      ...      ...      ...      ...      ...
839        1         1           3           1           0           1.000000
840        1         1           1           0           0           0.842199
841        1         1           1           0           0           0.000000
842        1         1           2           1           0           0.000000
843        1         0           0           0           0           0.000000
```

```
      Property_Area  ApplicantIncomeLog  LoanAmountLog  LoanAmount_Term_log  \
0                2          8.674197          4.859812          5.888878
1                0          8.430327          4.859812          5.888878
2                2          8.006701          4.204693          5.888878
3                2          7.857094          4.795791          5.888878
4                2          8.699681          4.955827          5.888878
..      ...      ...      ...      ...
839              2          8.292298          4.859812          5.198497
840              0          7.539559          4.127134          5.888878
841              0          7.933080          4.990433          5.888878
842              2          7.969012          3.828641          5.198497
843              1          8.334952          4.595120          5.888878
```

```
      Total_Applicant_Income_Log
0          8.674197
1          8.714732
2          8.006701
3          8.505525
4          8.699681
..      ...
839        8.292298
840        7.539559
```

```

841          8.456381
842          7.969012
843          8.334952

```

[844 rows x 11 columns]

```
[42]: Y_resampled
```

```

[42]: 0      1
      1      0
      2      1
      3      1
      4      1
      ..
      839    0
      840    0
      841    0
      842    0
      843    0
Name: Loan_Status, Length: 844, dtype: int32

```

```
[43]: df_resampled['Loan_status'].value_counts()
```

```

[43]: Loan_status
      1    422
      0    422
Name: count, dtype: int64

```

```

[44]: ## Splitting Train and Test Data
X_resampled_train, X_resampled_test, Y_resampled_train, Y_resampled_test = \
    train_test_split(X_resampled, Y_resampled, test_size = 0.25, random_state= 42)

```

```

[45]: ## First Model 1 : Logistic Regression with Balanced Dataset

model1_b = LogisticRegression()
model1_b.fit(X_resampled_train, Y_resampled_train)
y_pred_model1_b = model1_b.predict(X_resampled_test)
accuracy = accuracy_score(Y_resampled_test, y_pred_model1_b)
print("Accuracy Score of Logistic Regression: ", accuracy * 100)

```

Accuracy Score of Logistic Regression: 69.19431279620854

```

[46]: ## Second Model 2 : Decision Tree with Balanced Dataset

model2_b = DecisionTreeClassifier()
model2_b.fit(X_resampled_train, Y_resampled_train)
y_pred_model2_b = model2_b.predict(X_resampled_test)
accuracy = accuracy_score(Y_resampled_test, y_pred_model2_b)

```

```
print("Accuracy Score of Decison Tree : " ,accuracy * 100)
```

Accuracy Score of Decison Tree : 81.04265402843602

```
[47]: ## Third Model 3 : Random Classifier with Balanced Dataset
```

```
model3_b = RandomForestClassifier()
model3_b.fit(X_resampled_train,Y_resampled_train)
y_pred_model3_b = model3_b.predict(X_resampled_test)
accuracy = accuracy_score(y_pred_model3_b,Y_resampled_test)
print("Accuracy Score of Random Forest: " ,accuracy * 100)
```

Accuracy Score of Random Forest: 88.62559241706161

```
[48]: ## Fourth Model 4 : KNeighbor Model with Balanced Dataset
```

```
model4_b = KNeighborsClassifier(n_neighbors = 3)
model4_b.fit(X_resampled_train,Y_resampled_train)
y_pred_model4_b = model4_b.predict(X_resampled_test)
accuracy = accuracy_score(y_pred_model4_b,Y_resampled_test)
print("Accuracy Score of KNeighbor Model: " ,accuracy * 100)
```

Accuracy Score of KNeighbor Model: 72.51184834123224

```
[49]: ## Checking the reports of all models with balanced dataset
```

```
from sklearn.metrics import classification_report

def generate_classification_report(model_name, Y_test, y_pred):
    report = classification_report(Y_test, y_pred)
    print(f"Classification Report for {model_name}:\n{report}\n")

generate_classification_report(model1_b, Y_resampled_test, y_pred_model1_b)
generate_classification_report(model2_b, Y_resampled_test, y_pred_model2_b)
generate_classification_report(model3_b, Y_resampled_test, y_pred_model3_b)
generate_classification_report(model4_b, Y_resampled_test, y_pred_model4_b)
```

Classification Report for LogisticRegression():

	precision	recall	f1-score	support
0	0.84	0.54	0.65	114
1	0.62	0.88	0.72	97
accuracy			0.69	211
macro avg	0.73	0.71	0.69	211
weighted avg	0.73	0.69	0.69	211

Classification Report for DecisionTreeClassifier():

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.80	0.86	0.83	114
1	0.82	0.75	0.78	97
accuracy			0.81	211
macro avg	0.81	0.81	0.81	211
weighted avg	0.81	0.81	0.81	211

Classification Report for RandomForestClassifier():

	precision	recall	f1-score	support
0	0.92	0.86	0.89	114
1	0.85	0.92	0.88	97
accuracy			0.89	211
macro avg	0.89	0.89	0.89	211
weighted avg	0.89	0.89	0.89	211

Classification Report for KNeighborsClassifier(n_neighbors=3):

	precision	recall	f1-score	support
0	0.73	0.77	0.75	114
1	0.71	0.67	0.69	97
accuracy			0.73	211
macro avg	0.72	0.72	0.72	211
weighted avg	0.72	0.73	0.72	211

5 Predictions using selected model which is Random forest

```
[50]: # Prepare new data (matching the training features)
new_data = {
    'Gender': 0, # Example values
    'Married': 0,
    'Dependents': 0,
    'Education': 0,
    'Self_Employed': 0,
    'Credit_History': 1,
    'Property_Area': 2,
    'ApplicantIncomeLog': 10.82, # Log-transformed value
    'LoanAmountLog': 11.51, # Log-transformed value
    'LoanAmount_Term_log': 5.39, # Log-transformed value
    'Total_Applicant_Income_Log': 11.93, # Log-transformed value
}
```

```

}

# Create a DataFrame with the same structure as X_resampled
new_df = pd.DataFrame([new_data])

# Step 3: Make predictions
prediction = model3_b.predict(new_df)
if prediction[0] == 1:
    print("The person is eligible for a loan.")
else:
    print("The person is not eligible for a loan.")

# Step 4: (Optional) Get probabilities
probabilities = model3_b.predict_proba(new_df)
print(f"Probability of not eligible (Class 0): {probabilities[0][0]:.2f}")
print(f"Probability of eligible (Class 1): {probabilities[0][1]:.2f}")

```

```

The person is eligible for a loan.
Probability of not eligible (Class 0): 0.23
Probability of eligible (Class 1): 0.77

```

[]: